

SIH 2026 — Final Problem Statement Report for GlobeX

Six candidates, measured against what GlobeX actually is

Why each one fits | How we would build it | What makes us win

26132 26100 26102 26021 26002 26090

Prepared: 1 September 2026

All 229 published problem statements searched

Candidate selection was not done by keyword. A profile of GlobeX's verified architecture was scored against all 229 problem statements using TF-IDF cosine similarity — the same retrieval method GlobeX's own RAG engine uses — and the ranking is reported in Section 2.
One problem statement came out a clear outlier at the top.

1 The short version

What the search found

PS 26132 (Strengthening market linkages and price discovery for farmers) ranked **first out of all 229** problem statements, and not narrowly — its similarity score is **65% higher than the second-placed statement**. Reading it confirms why: it asks for market intelligence, price forecasting, verified buyer matching, lot creation, logistics coordination, payment tracking, dispute handling and transparent transaction records. That is GlobeX’s pipeline, moved from cross-border trade to domestic agriculture. It is the closest thing to “GlobeX, restated as a problem statement” in the entire dataset.

PS	Short title	Sim. rank	Coverage	The one-line case
26132	Market linkage & price discovery for farmers	1 / 229	9 of 9	Our product, restated for agriculture. Escrow answers a problem the PS names out loud.
26100	GeM bid compliance verification	4 / 229	14 of 14	Uses the most of GlobeX at once; the PS itself demands an audit trail.
26102	MPLAD anomaly & fraud detection	31 / 229	9 of 9	Our strongest trained model, transferred almost unchanged.
26021	Honey Chain traceability	30 / 229	7 of 8	Blockchain provenance at the centre, with escrow as the twist.
26002	NER smart logistics platform	40 / 229	6 of 8	Strong analytics fit; live GPS is new build.
26125	Blockchain identity, access & asset platform	1 / 229 (chain profile)	6 of 8	The pure blockchain twin — deepest chain fit in the dataset, but zero ML content.
26130	Industrial approvals & compliance single window	9 / 229 (compliance)	8 of 8	New find. 26100’s mirror image, from the applicant’s side.
26045	IP-SAKTI: multilingual source-cited RAG assistant	4 / 229 (compliance)	6 of 7	New find. The closest twin to our citation-enforced RAG engine.
26090	Artisan market linkage app	15 / 229	6 of 9	Marketplace core fits; generative AI and mobile are new.

2 Finding GlobeX’s twin among 229 problem statements

A profile of GlobeX’s verified capabilities — market discovery, demand and price forecasting, multi-criteria ranking, counterparty matching and trust scoring, risk and anomaly detection, compliance retrieval, document verification, blockchain evidence, escrow with disputes, logistics estimation — was vectorised alongside all 229 problem statements and scored by cosine similarity.

#	PS	Score	Title (short)
1	26132	0.076	Farmer market linkage & price discovery
2	26188	0.046	Fake identity & document screening
3	26006	0.042	Freight forecasting, vessel chartering
4	26100	0.041	GeM bid compliance verification
5	26089	0.037	Cooperative gig services platform
6	26229	0.034	Kabadiwala Connect recycling chain
7	26182	0.033	Crypto wallet attribution
8	26228	0.033	CV integrity assurance
9	26033	0.032	Farmer/consumer marketplace
10	26018	0.032	Land record digitisation
11	26108	0.032	Applicable-standards recommender
12	26099	0.031	Material code harmonisation
13	26046	0.031	Clinical trials dashboard
14	26024	0.031	Coal mine governance
15	26090	0.030	Artisan market linkage app

Reading the numbers

Absolute cosine values are low for everything — a short capability profile against long problem-statement prose produces sparse overlap. What matters is the *ordering* and the *gap*, and the gap at the top is decisive: 26132 sits 65% above second place, while ranks 2–15 are packed within a narrow band of each other.

Note also rank 9: PS 26033 is the thin, three-line version of the same farmer-marketplace idea that 26132 states properly. Where two statements describe the same problem, take the one with the detail — the rubric rewards root-cause understanding, and 26132 gives us material to work with.

Verdict on “is there a project like GlobeX in here?” Yes — one. PS 26132 is a genuine product-level twin. Everything below it matches on *technology* rather than on *product shape*.

Extended sweep: four profiles, not one

A single product-shaped query favours marketplace problem statements and can miss a statement that is a deep match to one *part* of GlobeX. The search was therefore re-run with four separate capability profiles across all 229 statements. Each column below is an independent ranking.

#	Product shape	Blockchain / trust	ML / anomaly / risk	Compliance / RAG
1	26132 farmer market linkage	26125 blockchain identity & assets	26170 component burn-in anomaly (ISRO)	26100 GeM bid compliance
2	26033 farmer market-place (thin)	26194 student innovation (open brief)	26070 AI/ML general	26107 BIS standards assistant
3	26089 cooperative gig platform	26211 student innovation (hardware)	26071 rainfall early warning	26188 fake identity screening
4	26229 Kabadiwala recycling	26021 Honey Chain	26017 land acquisition delay	26045 IP-SAKTI RAG assistant
5	26006 freight forecasting	26141 digital signature security	26103 project monitoring	26024 coal mine governance
6	26090 artisan linkage	26018 land record digitisation	26178 environmental monitoring	26190 legal document management
7	26136 startup procurement	26228 CV integrity assurance	26025 mine subsidence	26018 land record digitisation
8	26028 railway dynamic ETA	26046 clinical trials dashboard	26073 weather station anomaly	26035 test report generation
9	26164 cryptographic discovery	26190 legal document management	26102 MPLAD anomaly	26130 industrial approvals
10	26146 bitcoin traffic analysis	26100 GeM bid compliance	26183 crypto fraud exchanges	26036 metrology verification

What the extended sweep changed. Three things.

1. PS 26125 is the blockchain twin, and it is not close. It ranks **first of 229** on the chain/trust profile with a score roughly *70% above* second place — the same kind of decisive gap 26132 shows on product shape. Between them, the two statements bracket GlobeX: 26132 matches what the product *is*, 26125 matches what the chain layer *does*. Note the mirror-image weakness though: 26125 ranks **200th of 229** on the ML profile, because it contains essentially no machine learning at all.

2. Two genuinely new candidates surfaced — PS 26130 and PS 26045, both invisible to the original product-shaped query and both strong enough to earn a place in the comparison (Section 10).

3. Nothing better was found on the ML axis. The ML profile's top hits are real anomaly problems, but they sit in domains we have no data for and no product story in — component burn-in testing at ISRO, rainfall inundation, weather-station telemetry, mine subsidence. PS 26102 remains the best ML pick because it combines the same modelling shape with published data and a defensible blockchain angle.

3 Which parts of GlobeX each problem statement actually uses

● = core to the solution ○ = partially used ○ = not used ✕ = must be built new

	AI / ML	Deep learning	Forecasting	Anomaly / risk	NLP / RAG	OCR / docs	Blockchain	Escrow	Marketplace	Ranking / MCDA	Geo / GIS	Workflow	Dashboards
26132	●	○	●	○	○	○	●	●	●	●	○	●	●
26100	●	○	○	●	●	●	●	○	○	○	○	●	●
26102	●	●	○	●	○	○	●	○	○	○	○	●	●
26021	●	○	●	●	○	○	●	●	●	○	○	●	●
26002	●	●	●	●	○	○	○	○	○	●	●	●	●
26125	○	○	○	○	○	○	●	●	○	○	○	●	●
26130	●	○	○	○	●	●	○	○	○	○	○	●	●
26045	●	○	○	○	●	○	○	○	○	○	○	●	○
26090	●	○	●	○	●	○	○	○	●	○	○	●	●

What each one still needs built

PS	Net-new work
26132	✗ mandi price feed ingestion ✗ quality-grading capture
26100	<i>none</i> — every stated requirement lands on existing capability
26102	<i>none</i> — public scheme data is already published
26021	✗ IoT hive sensor ingestion ✗ farmer mobile app
26002	✗ live GPS telemetry ✗ offline sync ✗ multilingual notifications
26125	✗ DID / verifiable credentials ✗ NFT minting ✗ on-chain RBAC
26130	✗ department workflow connectors
26045	✗ IP/GI corpus ingestion ✗ multilingual answering
26090	✗ generative image enhancement ✗ voice-to-catalogue NLP ✗ native mobile app

4 PS 26132 — Market linkages and price discovery for farmers

Org: Government of Maharashtra (MSInS) | **Theme:** Agriculture, FoodTech & Rural Development | **Category:** Software | **Similarity rank:** 1 / 229 | **Coverage:** 9 of 9 | **Deadline:** 20 Sep 2026

Why this one

Smallholders and producer groups cannot see current or expected prices across nearby mandis, processors and institutional buyers. Information on quality specifications, demand, logistics, storage, payment reliability and buyer credentials is scattered. Farmers sell immediately at harvest because they lack storage and liquidity, and they negotiate from a weak position. Buyers, on the other side, cannot aggregate consistent volumes or verify quality.

That is an information-asymmetry problem wrapped around a trust problem — which is exactly the problem GlobeX was built to solve, one market up. An exporter who cannot see which country will pay best, cannot tell a reliable buyer from an unreliable one, and cannot be sure of getting paid, is the same person as a farmer at a mandi gate. We already built: price and demand forecasting to answer *when and where to sell*, a multi-criteria ranking engine to turn many signals into one recommendation, counterparty matching with trust scoring to answer *who to sell to*, escrow with milestone conditions and dispute handling to answer *will I actually get paid*, and a tamper-evident ledger to make the record of it all independently checkable.

The problem statement then asks, in its own words, for "transparent transaction records", "payment reliability", "buyer credentials", "dispute or grievance processes" and "sale-window recommendations". We have a working implementation of every one of those ideas. Nothing here has to be invented; it has to be re-pointed.

How we would build it

	PS asks for	GlobeX	How it is built
1	Aggregate mandi prices, buyer demand, arrival volumes	Ingestion + feature pipeline	Point the existing trade-panel ingestion at Ag-marknet/eNAM mandi feeds. Same shape of data: a price series per commodity per market per day, instead of per corridor per month.
2	Localised price trends and sale-window recommendations	Forecasting (F1/F2)	The dual-head forecaster already predicts demand <i>and</i> price together — which is precisely what a "hold or sell now" recommendation needs. Use the quantile forecaster for the confidence band around the advice.
3	Match farmers/FPOs with verified buyers	Matching + trust score	Semantic matching over commodity, grade, volume and location, then the composite trust score — built for exactly this "is this counterparty credible" question.
4	Quality specifications and grading	Document intelligence	Grade certificates and lot records go through the existing extraction and reconciliation path; grade claims are checked against the lot record rather than taken on trust.
5	Lot creation and digital offers	Marketplace	Listings, offers and catalogue pages already exist; the seller becomes an FPO and the unit becomes a lot.

	PS asks for	GlobeX	How it is built
6	Logistics and storage coordination	ETA + cost estimator	The distance/time and landed-cost model transfers directly; road transport constants replace ocean-freight ones.
7	Payment tracking	Escrow life-cycle	Buyer funds escrow on offer acceptance; release is conditioned on delivery and grading, exactly like the existing DOCS / SHIPMENT / INSPECTION conditions.
8	Dispute or grievance process	Dispute + resolve states	Already implemented as first-class states in the escrow contract, with the funds held rather than lost while a dispute is open.
9	Transparent transaction records	Evidence ledger	Every offer, acceptance, grading and settlement is hash-anchored, so the record a farmer sees is the record the buyer sees and neither can quietly edit it.

Our USP on 26132 — escrow answers a problem the PS names out loud

Most teams will build a price dashboard with a buyer directory. A price dashboard does not fix the farmer's actual fear, which the problem statement states plainly: **payment reliability**. Our differentiator is that the transaction is protected, not just informed — the buyer's money is committed into escrow before the lot moves, release is conditioned on delivery and grading, a dispute freezes the funds instead of losing them, and the whole sequence is anchored so it can be audited later. That converts "here is a better price signal" into "here is a trade you can safely complete", which is a categorically stronger answer to the same problem statement. It is also the hardest part for another team to build in a hackathon window — and we already have it running.

Risks to manage

Mandi price data is public but messy — inconsistent commodity naming across markets, missing days, and units that vary by state; budget real time for the ingestion clean-up. Avoid over-claiming on the forecaster: use the quantile model and present a range with a confidence band rather than a single number, and never show a sale-window recommendation without the uncertainty around it.

5 PS 26100 — AI bid compliance verification for GeM procurement

Org: Ministry of Petroleum & Natural Gas (CPCL) | **Theme:** Smart Automation | **Category:** Software | **Similarity rank:** 4 / 229 | **Coverage:** 14 of 14 | **Deadline:** 20 Sep 2026

Why this one

Procurement officers verify every bidder by hand across a dozen government registries — Udyam, GSTN, PAN, EPFO/ESIC, DigiLocker, debarment lists — then decide whether the bid qualifies. It is slow, inconsistent, and error-prone, and the consequence of getting it wrong is a contract awarded to an ineligible vendor.

Strip the domain away and this is the problem GlobeX solves before every trade: take an entity, check it against many authorities, reconcile the documents it submitted, score what came back, and give a human a defensible recommendation with the evidence attached. All fourteen numbered requirements in the problem statement land on something we have already built — there is no net-new capability in it at all, which is true of no other candidate except 26102.

How we would build it

	PS asks for	GlobeX	How it is built
1	Registry verification: Udyam, GSTN, PAN, EPFO, NSIC, Startup India	Entity screening	The screening engine already resolves an entity against a 31,629-record registry with fuzzy name matching; the registries swap, the matching logic does not.
2	Blacklisting / debarment status	Denied-party screening	A debarment list <i>is</i> a denied-party list. This is the same check, unmodified.
3	Make in India / local content	Rules-of-origin logic	Local-content percentage thresholds are computationally identical to value-addition rules of origin, which the compliance store already encodes.
4	Document verification and DigiLocker checks	OCR + reconciliation	Field extraction and cross-document consistency checking, already running over trade documents.
5	Identify missing or inconsistent information	Reconciliation	The routine that catches an invoice quantity disagreeing with a bill of lading catches a turnover figure disagreeing with a filed return.
6	Compliance score and bidder risk level	Composite scoring	Multi-factor scoring with banding, blending registry results and behavioural signals.
7	AI recommendation, officer keeps the decision	Transaction gate + dossier	A human-in-the-loop decision gate with a machine recommendation and a written justification — already how trade approval works.
8	Auditable record of every check	Evidence ledger	Each verification is hash-anchored with its result and timestamp.

Our USP on 26100 — an audit trail that cannot be edited after the fact

Requirement 14 asks for an auditable record. Most teams will satisfy that with a database log — which the awarding authority can edit. Ours is cryptographically anchored, so a tender decision made in March can be proved in November to have been based on exactly the evidence shown at the time. In public procurement, where the dispute is usually *"what did you actually check before you awarded this?"*, that is the difference between a log and proof. This is the one candidate where blockchain is demanded by the problem statement rather than added by us, so it cannot be dismissed as decoration.

Risks to manage

Live portal credentials (GSTN, DigiLocker) will not be issued to a student team — demo against the dummy bidder and tender datasets the problem statement itself offers, and draw the integration boundary explicitly in the architecture slide rather than implying live access.

6 PS 26102 — Anomaly, fraud and inefficiency detection in MPLADS

Org: MoSPI — Data Informatics & Innovation Division | **Theme:** Smart Automation | **Category:** Software | **Similarity rank:** 31 / 229 | **Coverage:** 9 of 9 | **Deadline:** 20 Sep 2026

Why this one

Thousands of local works are funded through many implementing agencies, producing large volumes of expenditure and execution data that nobody can review manually. The problem statement asks for ML that surfaces anomalies in expenditure, fund utilisation, cost estimates and work execution, so irregularity is caught early.

This is the closest match in the dataset to our strongest asset. The trade anomaly classifier is trained, chronologically validated and performing at 0.98 F1, and the unsupervised risk ensemble exists precisely because trade data has *no labelled fraud* — the same condition public-fund data is in. Most teams attacking this PS will have to invent an approach for unlabelled irregularity detection during the hackathon. We already have one, with attribution behind every flag.

How we would build it

	PS asks for	GlobeX	How it is built
1	Anomalies in expenditure patterns	XGBoost anomaly + features	Sanctioned-work expenditure records are the same object as trade transactions: an amount, a date, a category, a counterparty. The rolling-statistics and growth features transfer directly.
2	Irregularity without labelled fraud	Isolation Forest + autoencoder	The unsupervised ensemble runs unchanged. Reconstruction error over a work's spending sequence flags what does not resemble normal execution.
3	Cost-estimate deviation	Unit-value deviation	Our peer-price deviation feature becomes cost-per-unit against comparable works in comparable districts.
4	Duplicate or overlapping works	TF-IDF similarity	Near-duplicate work descriptions across constituencies are caught the same way near-equivalent company records are.
5	Risk scoring and prioritisation	Composite score	Ranks works and agencies so limited audit capacity goes where the signal is strongest.
6	Explainable flags	TreeSHAP attribution	Every flag carries its contributing factors — non-negotiable when the subject is public money.
7	Monitoring dashboard and alerts	Dashboards + n8n	Existing charting stack and resilient workflow orchestration.

Our USP on 26102 — blockchain as the accountability layer, not decoration

Here is the move most teams will miss. An anomaly-detection tool that accuses an implementing agency creates a new problem: who flagged this, on what evidence, and did anyone quietly change it afterwards? We anchor **the flag itself** — the model version, the input snapshot hash, the score and the officer's disposition — to the evidence ledger. That makes the oversight trail as tamper-evident as the spending it oversees, protects agencies from silently altered accusations, and protects auditors from claims of tampering.

Paired with our second differentiator — we do **not** claim the model detects fraud. Our own risk documentation refuses to equate an anomaly with fraud because there is no ground truth for it, so the output is "this pattern is unusual and warrants review", feeding a human investigation. In a domain where a false public accusation is a real harm, that restraint plus an immutable audit trail is a far stronger pitch than a higher accuracy number.

Risks to manage

Do not present the 0.98 F1 from trade data as a result on MPLADS data — it is evidence the method works, not a claim about this domain, and a judge who catches that conflation will discount everything else. Retrain and report honestly on the scheme's own published data.

7 PS 26021 — Honey Chain: blockchain traceability and smart beekeeping

Org: Ministry of MSME (KVIC Honey Mission) | **Theme:** Agriculture, FoodTech & Rural Development | **Category:** Software | **Similarity rank:** 30 / 229 | **Coverage:** 7 of 8 | **Deadline:** 20 Sep 2026

Why this one

KVIC-supported beekeepers face counterfeit honey diluting their market, weak buyer linkage, and no way to prove their product is what they say it is. The problem statement asks for blockchain-based batch traceability with QR consumer verification, plus AI/IoT support for disease detection and yield prediction.

Structurally this is GlobeX's spine with a different commodity: producers listed in a marketplace, provenance anchored on-chain, forecasting on the supply side, anomaly detection on the operational side. The chain layer here is not bolted on — authenticity *is* the problem, which makes it the most natural home for our blockchain work of any candidate except 26100.

How we would build it

	PS asks for	GlobeX	How it is built
1	Batch traceability, hive to shelf	Chain anchor + ledger	Each batch event (harvest, extraction, packing, dispatch) is hash-anchored using the existing write-intent-then-confirm pattern — the payload changes, the guarantee does not.
2	QR consumer verification	Hash re-verification	Scanning re-computes the hash and compares it against the anchored record: our document integrity check, exposed through a QR code.
3	Market linkage for producers	Marketplace + matching	Listings and buyer matching already exist; sellers become beekeeper collectives.
4	Productivity / yield prediction	Quantile forecaster	Per-cluster yield forecasting with confidence bands. Use the quantile model, not the GRU.
5	Disease and environmental monitoring	Anomaly detection	Anomaly detection over hive sensor series is the same problem as over trade series; only the ingestion path is new.
6	Authenticity and market credibility	Escrow	See USP below.

Our USP on 26021 — payment released against verified quality

"Blockchain + AI + QR for traceability" is the most crowded pattern in any hackathon; assume several teams present it. What they will not have is the second half of the trust problem. Traceability proves where honey came from but does nothing about whether the beekeeper actually gets paid fairly for it. We add **quality-conditioned escrow**: the buyer commits funds when the batch is listed, release is conditioned on the inspection result recorded against that batch, and a dispute holds the money rather than leaving the producer to chase it. Provenance plus guaranteed settlement is a complete answer to "why does counterfeit honey persist" — provenance alone is only half of it.

Risks to manage

There is no hive sensor dataset available — the IoT layer must be demonstrated on simulated or public apiary data, stated openly. And because the pattern is crowded, the pitch must lead with root cause (why existing certification has failed to stop dilution) rather than with the technology stack.

8 PS 26125 — Blockchain identity, access control and digital asset management

Org: Bharat Electronics Limited | **Theme:** Blockchain & Cybersecurity | **Category:** Software | **Similarity rank:** 1 / 229 on the chain profile | **Coverage:** 6 of 8 | **Deadline:** 20 Sep 2026

Why this one

Centralised identity and access management is a single point of failure: vulnerable to breach, identity theft and unauthorised access, and it leaves asset ownership scattered across disconnected systems where authenticity and ownership history cannot be reliably verified. The problem statement asks for a decentralised alternative — self-sovereign identifiers authenticated by cryptographic proof, digital assets issued as NFTs bound permanently to those identities, smart contracts enforcing who may mint, allocate and transfer, role-based access control enforced on-chain rather than in application code, and every operation immutably recorded as an audit trail.

This is the deepest chain-layer match in the dataset — first of 229 on the blockchain profile, roughly 70% clear of second place. The reason is that GlobeX’s escrow contract already embodies the exact discipline this statement is asking for: state transitions governed by a contract rather than by application code, actions permitted only to authorised parties, and a rule that no consequential state is ever written without a confirmed on-chain receipt behind it. That last point is the difference between a system that *uses* a blockchain and one that can actually be *trusted* because of it, and we already enforce it in production code.

The honest counterweight: this problem statement contains **almost no machine learning** — it ranks 200th of 229 on the ML profile. Choosing it means our forecasting, anomaly detection, risk ensemble and ranking work — the majority of the engineering behind GlobeX — does not appear in the solution at all. It is the exact mirror image of PS 26102’s trade-off.

How we would build it

	PS asks for	GlobeX	How it is built
1	Smart contracts governing all operations	Escrow contract + adapter	We already run a lifecycle contract with guarded state transitions through a dedicated chain-adapter service; the contract’s subject changes from a trade to an identity and its assets.
2	Immutable audit trail of every operation	Evidence ledger	Identity creation, minting, allocation, transfer and permission changes are anchored exactly as document and trade events are anchored today.
3	Ownership and authenticity verification	Hash verification	Re-deriving a hash and comparing it against the anchored record is our existing integrity check.
4	Controlled issuance — only authorised admins may mint	Role-gated operations	The escrow contract already restricts who may fund, release and resolve; the same authorisation pattern governs minting.
5	Ownership history	Ledger query + UI	The blockchain ledger screen already renders anchored event history for a trade; here it renders it for an asset.
6	Transaction confirmation discipline	Write-intent pattern	Our submit-then-confirm rule (never record a final state on an unconfirmed transaction) transfers wholesale, and is the single most valuable thing we bring.

	PS asks for	GlobeX	How it is built
7	Decentralised identifiers & verifiable credentials	✗ new	DID issuance and cryptographic proof verification is new contract and client work.
8	NFT minting and on-chain RBAC	✗ new	Token standard implementation and on-chain role enforcement go beyond our current escrow contract.

Coverage: 6 of 8 asks land on existing capability; 2 require new Solidity work.

Our USP on 26125 — the discipline, not the buzzwords

Every team attempting this problem statement will demo a contract that mints an NFT. Most will also, somewhere in the stack, write "transferred" into a database before the chain has confirmed it — because that is the shortcut that makes a demo work on a flaky testnet. Our differentiator is the failure handling nobody else will have thought about: we already write a pending intent record *before* the chain call, patch it to confirmed with the real transaction hash on success or to failed with the real error on failure, and never let the application claim an ownership change the chain has not acknowledged. In a problem statement whose entire premise is trustworthiness, being the one team whose system cannot lie about its own state is a far stronger position than having the fanciest token design.

The trade-off to weigh

This is the narrowest of the nine candidates. It exercises our chain layer more deeply than anything else, and nothing else at all — no forecasting, no anomaly detection, no risk scoring, no ranking, no compliance retrieval. If the team's judgement is that GlobeX's identity is the blockchain work, this is the strongest possible showcase for it. If the ML work is the identity, this is the weakest of the nine.

9 PS 26002 — AI smart logistics and accessibility platform (NER)

Org: Ministry of Development of North Eastern Region | **Theme:** Transportation & Logistics | **Category:** Software
 | **Similarity rank:** 40 / 229 | **Coverage:** 6 of 8 | **Deadline:** 20 Sep 2026

Why this one

Landslides, floods and fragile terrain cut road links to remote NER districts, delaying medicines, food and produce. The problem statement asks for a platform that monitors accessibility, predicts disruptions, recommends alternate routes, tracks vehicles, alerts authorities, accepts geo-tagged field reports and presents district-level dashboards.

The analytical core maps well, and one mapping is genuinely elegant: **ranking alternate routes is the same problem as ranking destination markets**. Our multi-criteria engine already turns eight weighted, explainable signals into one ranked recommendation — point it at route options scored on rainfall risk, slope, incident history, detour cost and time, and it does the same job. Corridor-level risk scoring is also already implemented; today the corridors are trade lanes, here they are road segments.

How we would build it

	PS asks for	GlobeX	How it is built
1	Predict route disruptions	Forecast + anomaly risk	The same three-model shape we run on trade corridors, with rainfall, slope and incident history replacing trade features.
2	Alternate route ranking + travel delay	MCDA ranking + ETA	Weighted, explainable route scoring; delay estimated by the distance-over-speed method with terrain and condition penalties.
3	High-risk corridor identification	Risk ensemble	Corridor-level risk scoring, already implemented — the object the PS calls a "high-risk transport corridor" is the object K1 already produces.
4	Automated alerts	n8n workflows	Resilient multi-step alerting with per-step status, already built.
5	District connectivity dashboards	Dashboards + geo view	Existing charting and geospatial visualisation stack.
6	Geo-tagged field reports	Document intake	Upload, extraction and persistence exist; geo-tagging and image handling are additions.
7	GPS vehicle tracking	✗ new	Live telemetry ingestion, position storage and streaming. Our ETA service is a static estimator, not tracking.
8	Offline sync + multilingual alerts	✗ new	Offline-first synchronisation for low-network areas is a discipline we have not built.

Our USP on 26002 — explainable route decisions, not a black box

Most teams will produce a map with a shortest-path line on it. Ours ranks routes on weighted criteria and *shows why* — this route scores lower because of slope-adjusted landslide risk and recent incident density, not merely because it is longer. For a district officer deciding whether to send a medicine convoy, an explainable, auditable recommendation is the difference between a tool they trust and one they override. Our backtested ranking methodology is the credible core here.

Be honest about two things

Live GPS tracking is a headline feature of this platform and we would be building it from zero — say so rather than implying we have it. And the chain and escrow half of GlobeX has no natural role here; forcing it in reads as technology for its own sake, which the first-round rubric penalises. This is the one candidate where our blockchain work is a spectator.

10 PS 26090 — AI market linkage and smart cataloging for artisans

Org: Ministry of Social Justice and Empowerment | **Theme:** Heritage & Culture | **Category:** Software | **Similarity rank:** 15 / 229 | **Coverage:** 6 of 9 | **Deadline:** 20 Sep 2026

Why this one

Marginalised artisans and weavers get a sales bump from physical exhibitions but have no year-round digital channel, blocked by low digital literacy, language barriers and the practical difficulty of photographing, pricing and cataloguing products to e-commerce standards. The ask is a mobile app acting as a "virtual business manager": AI image enhancement, voice-to-catalogue in regional languages, and a dynamic pricing assistant, connecting artisans to B2B buyers and government e-marketplaces.

The commercial half of this is our marketplace: catalogue and listing creation, buyer discovery and matching, and price recommendation — all of which exist, including the listing and catalogue editor screens. The reason it sits below the others is the other half: generative image enhancement, speech-to-text in regional languages and generated product copy are capabilities we have not built, and the deliverable is explicitly a native mobile app rather than a web platform.

How we would build it

	PS asks for	GlobeX	How it is built
1	Dynamic pricing assistant	Forecasting + ranking	Price recommendation from comparable-product signals with a confidence band — our forecasting and scoring path, applied to craft categories.
2	Connect artisans to B2B buyers / GeM	Marketplace + matching	Listings, buyer discovery, semantic matching and trust scoring already exist.
3	Product cataloguing and inventory	Catalogue editor	Catalogue and listing creation screens are already built.
4	Buyer credibility for vulnerable sellers	Trust scoring + escrow	See USP below.
5	AI image enhancement studio	✗ new	Background removal and lighting correction — a generative/CV capability we do not have.
6	Voice-to-catalogue in regional languages	✗ new	Speech recognition plus generated multilingual product copy.
7	Native cross-platform mobile app	✗ new	Our frontend is web; this PS explicitly wants mobile-first for low-literacy users.

Our USP on 26090 — protecting the seller, not just listing them

Every team will build the camera and the auto-description; those features are the obvious reading of this problem statement. The risk nobody addresses is that a low-literacy artisan, newly exposed to B2B buyers, is the party most vulnerable to non-payment and exploitative terms. Our differentiator is **verified buyers plus protected payment** — trust-scored buyer credentials and escrowed settlement so the artisan ships against committed funds rather than a promise. For a Ministry of Social Justice problem statement, protecting the beneficiary from the market they are being introduced to is a substantially more thoughtful answer than a prettier product photo.

Why it is ranked sixth

Three of the nine asks are net-new capability, and two of them (generative imaging, multilingual speech) are the features the problem statement leads with. We would be strongest on the half of the problem the judges may treat as secondary — a real risk of a mismatch between our strength and the rubric's weight.

11 New finds from the extended sweep

Neither of these appeared in the original product-shaped search. Both were surfaced by the compliance/RAG profile, and both are strong enough to sit in the main comparison. The problem-statement Excel was checked row by row against the searched dataset first — 229 IDs, no differences in title, description, category or theme — so this sweep covers the complete published set.

PS 26130 — Streamlining industrial approvals, compliance and support services

Org: Government of Maharashtra (MSInS) | **Theme:** Miscellaneous | **Category:** Software | **Similarity rank:** 9 / 229 on the compliance profile | **Coverage:** 8 of 8 | **Deadline:** 20 Sep 2026

Why. Entrepreneurs setting up an industrial unit must obtain registrations, licences, NOCs, inspections and renewals from many authorities, and the requirements vary by sector, location, project size and stage of operation. Applicants cannot work out which approvals apply to them; departments receive incomplete applications and cannot see where the bottlenecks are.

This is PS 26100 seen from the other side of the counter. Where 26100 asks us to verify a bidder against authorities, 26130 asks us to tell an applicant which authorities apply and walk them through it — and the core of that request, “*generate a customised approval checklist*”, is precisely what our required-document matrix does today: take an item’s classification plus its origin and destination, and return the exact set of documents and clearances that apply to it. Swap HS6 code and trade lane for sector and location and the function is unchanged.

How. The applicable-approval checklist comes from the document matrix; documentation guidance comes from the citation-backed retrieval engine; pre-validation of submissions reuses the extraction and inconsistency-detection path; “reuse verified data” is our verified-entity record; parallel departmental coordination, inspection scheduling and SLA timers are n8n workflows with per-step status already built for exactly that failure-isolation reason; bottleneck visibility is the dashboard layer. All eight stated outcomes land on existing capability — only the department-side connectors are new.

USP on 26130 — the checklist is the hard part, and we already built it

Most teams will build a portal: forms, uploads, a status tracker. The genuinely difficult piece is the one the problem statement leads with — correctly deriving *which* approvals apply to a given unit, when the answer depends on four interacting variables. That is a rules-and-retrieval problem we have already solved once for trade compliance. Adding chain-anchored approval records on top means an issued clearance cannot be quietly back-dated or altered, which matters in exactly the environment where approvals are contested.

PS 26045 — IP-SAKTI Sahayak: multilingual, source-cited RAG assistant

Org: Ministry of Ayush | **Theme:** MedTech / BioTech / HealthTech | **Category:** Software | **Similarity rank:** 4 / 229 on the compliance profile | **Coverage:** 6 of 7 | **Deadline:** 20 Sep 2026

Why. Protecting and commercialising an Ayurvedic product means navigating several overlapping legal regimes simultaneously — patents, geographical indications, trademarks, copyright, designs, trade secrets, plant-variety rights, access-and-benefit-sharing duties over biological resources, and a drug-regulatory framework that decides whether a formulation counts as a classical medicine, a proprietary medicine, a new drug, a phytopharmaceutical, a food or a cosmetic. The ask is a multilingual assistant that answers in plain language **with citations to the specific source**.

That is a closer description of our RAG engine’s actual job than any other statement in the dataset, including PS 26107. Our retriever already spans multiple overlapping regulatory corpora at once — tariff schedules, a sanctions registry, export-control lists, rules of origin, SPS/TBT standards — and answers against them with provenance attached. The classification question is familiar too: deciding

whether a formulation is a classical medicine or a phytopharmaceutical, and therefore which pathway applies, is the same shape as deciding an HS6 code and therefore which document set applies.

How. Retrieval, ranking and citation enforcement transfer directly. The versioned current-facts store matters more here than anywhere else, because this domain moved recently — the 2024 patent and biodiversity rules and the 2024 WIPO treaty — and an assistant citing superseded law is worse than no assistant. New work: ingesting the IP, GI and AYUSH regulatory corpora, and answering in multiple languages.

USP on 26045 — we already refuse to answer without a source

The obvious build here is a chatbot over some documents, and the obvious failure is a confident answer with an invented citation — which, in a legal and regulatory domain, is not a glitch but a liability for the person who acted on it. Our engine is built the other way round: provenance is mandatory and fabricated citations are prohibited by design. Leading with *refusal behaviour* — what the assistant does when it does not have a source, and how it shows which clause an answer rests on — is a credibility argument almost no competing team will make, and it is the argument that matters most to a ministry putting its name on legal guidance.

12 Recommendation

PS	Best when the team wants to lead with	Net-new to work	Blockchain role	Headline USP
26132	The whole product	Low	Central	Escrow answers "payment reliability", named in the PS
26100	Compliance & verification	None	Demanded by PS	Audit trail that cannot be edited afterwards
26102	The trained ML core	None	Added by us	Tamper-evident record of the flags themselves
26021	Blockchain & traceability	Medium	Central	Payment released against verified quality
26002	Analytics & geospatial	High	None honest	Explainable route ranking, not a black box
26125	Blockchain, purely	Medium	Everything	Escrow-grade contract discipline: no state written without a confirmed receipt
26130	Compliance, applicant side	Low	Optional	Approval checklist generation — our required-document matrix, re-pointed
26045	RAG with citations	Medium	None	Refusing to answer without a source — already enforced in our engine
26090	Marketplace & social impact	High	Light	Verified buyers + protected payment for artisans

The recommendation

Lead with PS 26132. It is the measured closest match to GlobeX in the entire dataset by a wide margin, every one of its nine asks lands on something we have built, and the differentiator is not a technology we are showing off but the direct answer to a problem the statement itself raises — farmers selling into a market where payment is not reliable. It also lets us present the full platform story (intelligence → matching → transaction → settlement → dispute) as one coherent solution rather than a collection of features.

Keep PS 26100 as the strong second. It is the only other candidate with zero net-new work, and it is the one where blockchain is required by the problem statement rather than offered by us — the safest possible ground for defending that part of the stack under questioning.

PS 26102 is the pick if the team would rather lead with the models, and the blockchain angle (anchoring the flags themselves) turns what looked like an unused half of GlobeX into that submission's sharpest idea.

PS 26002 remains viable and the route-ranking insight is real, but it carries the most net-new work of the six and leaves our most distinctive asset on the bench.

Where the three additions land. **PS 26125** is the strongest showcase for the chain layer specifically — first of 229 on that profile — but it is the narrowest option on the table, using none of the ML work.

PS 26130 is the most interesting of the two new finds: it matches 26100’s low-drift profile from the applicant’s side, all eight of its outcomes land on existing capability, and its hardest requirement is a function we have already built. **PS 26045** is the truest match to the RAG engine and displaces 26107 as the retrieval-led option, though its corpus ingestion is real work. None of the three displaces 26132 at the top.

One practical note. 26132, 26021 and 26090 all exercise the same marketplace, matching, escrow and ledger path, so preparation for any one transfers substantially to the others. 26100, 26130 and 26045 share the screening, retrieval, document and evidence path — that cluster is now three statements deep, which makes it the safest area to invest preparation in. 26102 shares the scoring half of it. The two outliers are 26002 (its telemetry work transfers nowhere) and 26125 (its contract work transfers only to 26021), so both should be decided on early rather than prepared in parallel.

A note on the search itself. The problem-statement Excel supplied by the team was verified against the dataset used here before this sweep was run: 229 rows, 229 unique IDs, and no difference in any title, description, category or theme. Every candidate in this report was therefore drawn from the complete published set, scored under four independent capability profiles.

Unrelated to PS choice, but worth fixing before any submission

The repository’s new `README.md` advertises a demand-forecast error of **MAPE 8.42%** (“Deep GRU + Quantile Loss”). The benchmark artifact still sitting in both model folders, `benchmark_comparison.csv`, reports the Dual-Head GRU at **54.49% MAPE, R² 0.399** — worse than a three-year moving average on the same data. Either a newer quantile model exists and the benchmark file is stale, or the headline figure is not backed by the artifact. Whichever it is, the two should be reconciled before that number appears in a submission, because it is the first thing a technical judge would check.

Candidates selected by TF-IDF cosine similarity of a GlobeX capability profile against all 229 problem statements from <https://www.sih.gov.in/sih2026PS>, then mapped against verified backend modules, trained model artifacts and frontend screens. GlobeX is an in-flight build — this maps capability direction, not final code. All six share a 20 September 2026 idea-submission deadline.